

High-Performance Graph & Data Analytics

Spring 2022

An efficient representation for sparse graphs:

Compressed Sparse Row

Bosisio Andrea – Karra Salvatore Gabriele



Goal

Implementation
of a graph data
structure for
sparse graphs

- Efficiently population:
`populate(edge_list)`
- Fast traversing:
`get_neighbors(vertex)`
- Limited memory usage

Why CSR?

Brief comparison with other data structures

- **Adjacency matrix**
 - Too much memory: $O(|V|^2)$
- **Adjacency List**
 - It was the example implementation
 - Adding a neighbor is not efficient (dynamic memory allocation)
- **COO**
 - A lot repeated information
 - Neighbors iteration: $O(|E|)$: we don't know where each vertex starts

Why CSR?

Advantages of using this data structure

- **Smart information storage**
 - No duplicated data
 - It exploits the fact that data are sorted
 - good space complexity (see later...)
- **Fast out-neighbors iteration**
 - Direct access to the first neighbor
 - Then iteration is linear:
 - If a vertex has N_n neighbors → $O(N_n)$
 - In general → $O(|V|)$

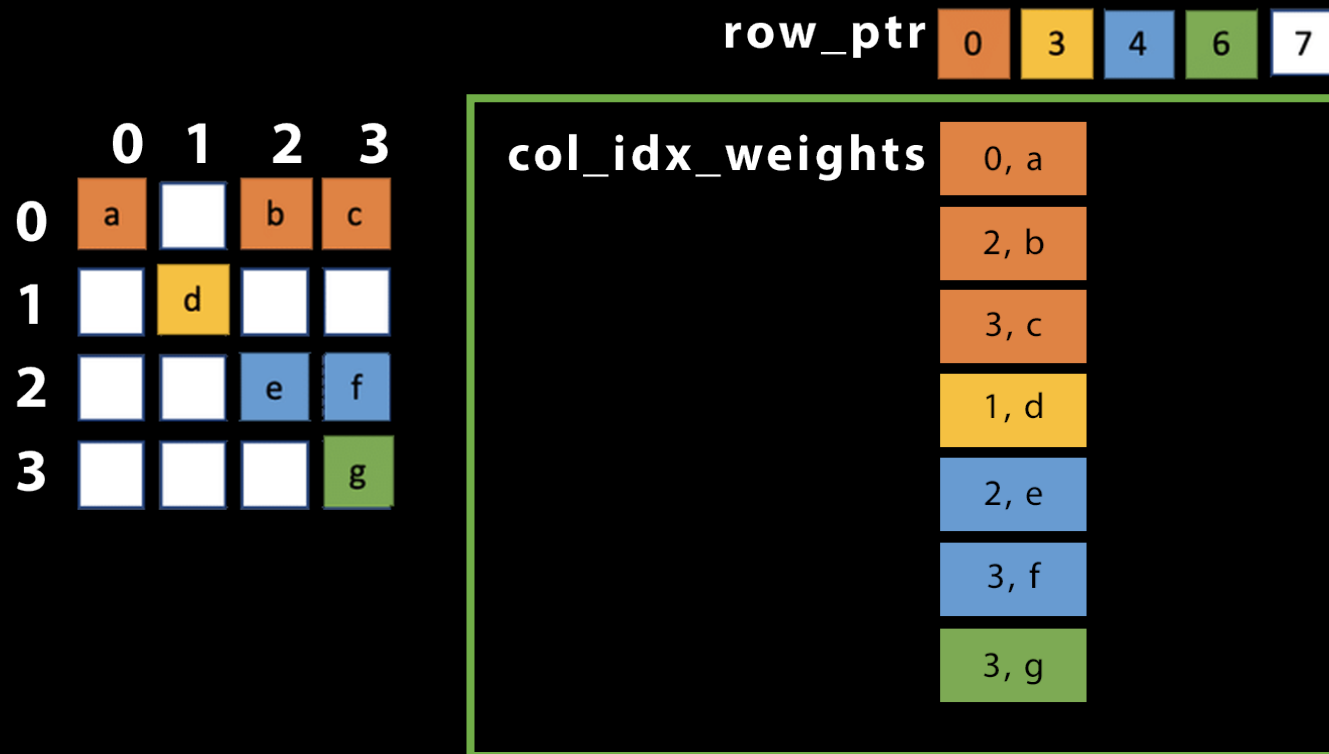
How does CSR work?

	0	1	2	3
0	a		b	c
1		d		
2			e	f
3				g

What we need to store?

	0	1	2	3	4		
row_ptr	0	3	4	6	7		$O(V)$
col_idx	0	2	3	1	2	3	$O(E)$
weights/value	a	b	c	d	e	f	$O(E)$
							TOT: $O(V + E)$

A first optimization



Why?

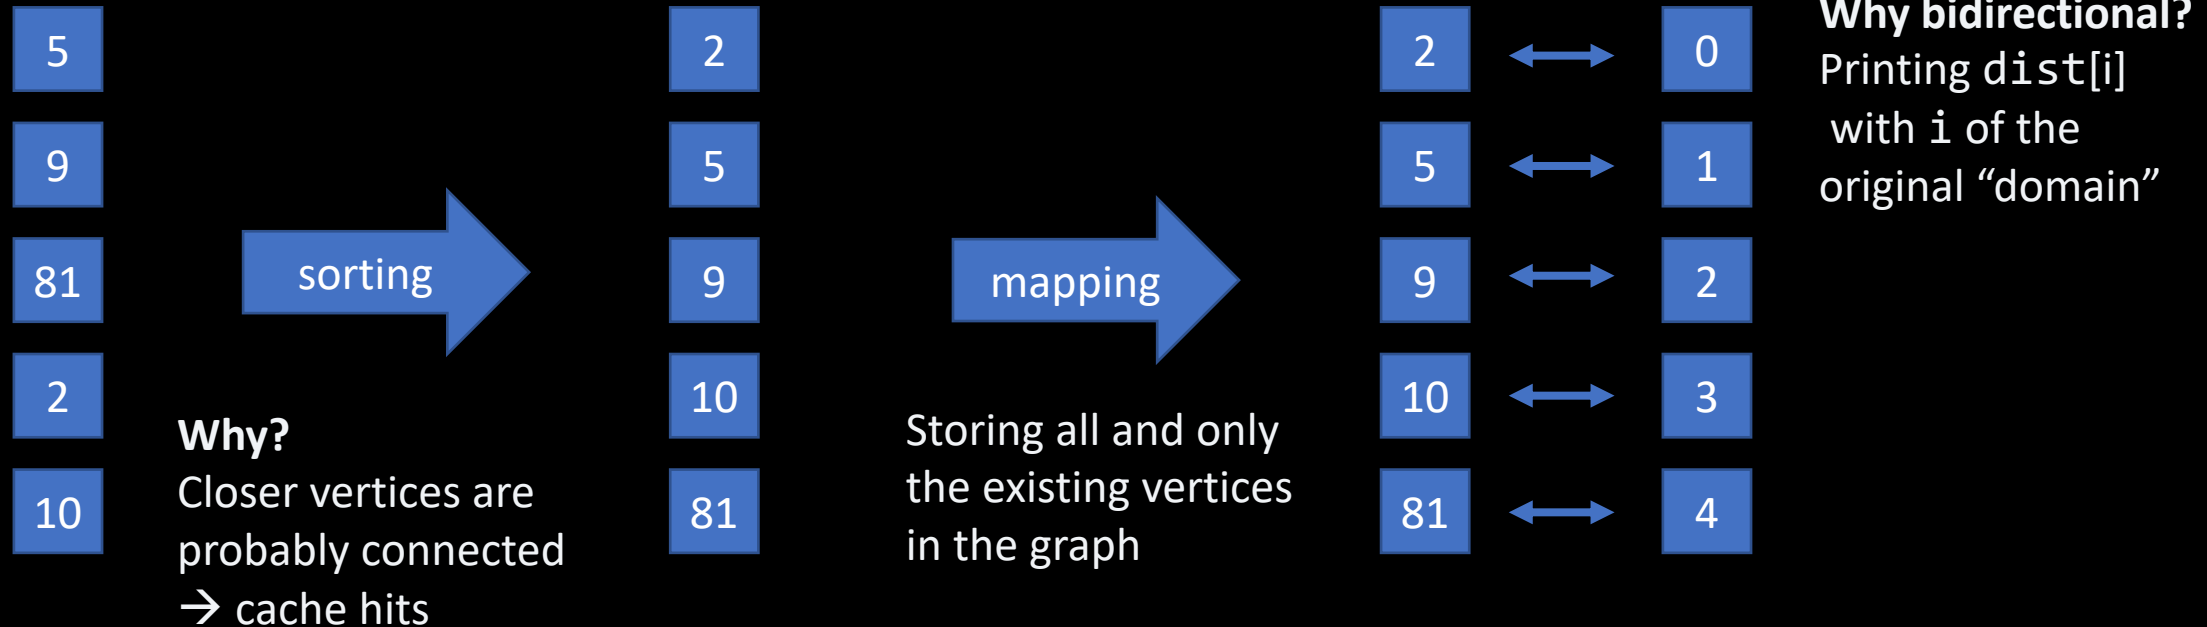
col_idx and **weights** are always accessed together:

- Reducing conflicts
- Improve spatial locality (operate on data in the same cache block)

→ Introducing **col_idx_weights** pair

Data pre-processing in `load_graph()`

Handling graphs with no sequential vertex indexes and
“isolated” vertices



Populating the CSR first problem

The edge list in input is NOT sorted and the CSR assumes that edges are sorted row by row from an adjacency matrix

We NEED to sort the edge list w.r.t. the first element of the tuple `<node_from, node_to, weight>`

We need to find a proper sorting algorithm...

Populating the CSR

a smart solution


row_ptr

0	3	4	6	7
---	---	---	---	---

This looks like something we have already seen in a famous sorting algorithm...

The **count** array of the *counting sort* does the same thing!

Counting Sort



We exploit this algorithm to:

- Find the `row_ptr` array that is basically the count array
- Building the `col_idx_weights` array which is now sorted w.r.t. the starting vertex (i.e. `node_from`)

Let's now see in detail the code...

populate()

#1. Counting the number of neighbors for each vertex

PARALLEL VERSION

```
#pragma omp parallel for reduction  
(+:row_ptr[: (num_vertices+2)])  
  
for(uint64_t i = 0; i < num_edges; i++)  
    row_ptr[std::get<0>(e_list[i]) + 1]++;
```

Time Complexity: $O\left(\frac{|E|}{P} + P\right)$
where P is the number of partitions
(i.e. threads)

SERIAL VERSION

```
for (uint64_t n = 0; n < num_edges; ++n)  
    row_ptr[std::get<0>(e_list[n]) + 1]++;
```

Time Complexity: $O(|E|)$

populate()

#2. Computing the cumulative sum & Copying row_ptr into count

```
// cumulative sum
std::partial_sum(row_ptr, row_ptr + (num_vertices+2), row_ptr);

// copying row_ptr into count which will be used for the counting sort
std::copy(row_ptr + 1, row_ptr + 1 + num_vertices, count);
```

Time Complexity: $O(|V| + |V|) = O(|V|)$

Drawback: we need to duplicate data on the heap
since count will be decremented

populate()

#3. Sorting and filling the col_idx_weight array

PARALLEL VERSION

```
initLocks();
#pragma omp parallel {
std::tuple<uint64_t, uint64_t, double> curr_edge;
uint64_t curr_vertex, new_pos;

#pragma omp for
for (uint64_t i = 0; i < num_edges; i++) {
    curr_edge = e_list[i]; curr_vertex = std::get<0>(curr_edge);
    // locking on curr_vertex
    omp_set_lock(&count_locks[curr_vertex]);
    new_pos = count[curr_vertex] - 1; count[curr_vertex]--;
    // unlocking curr_vertex
    omp_unset_lock(&count_locks[curr_vertex]);

    col_idx_weight[new_pos] =
        std::make_pair(std::get<1>(curr_edge),
std::get<2>(curr_edge));
}}

// continuing ...
```

SERIAL VERSION

```
std::tuple<uint64_t, uint64_t, double> curr_edge;
uint64_t curr_vertex, new_pos;

for (uint64_t i = 0; i < num_edges; i++)
{
    curr_edge = e_list[i];
    curr_vertex = std::get<0>(curr_edge);

    new_pos = count[curr_vertex] - 1;
    count[curr_vertex]--;

    col_idx_weight[new_pos] =
        std::make_pair(std::get<1>(curr_edge),
std::get<2>(curr_edge));
}
```

Time Complexity: $O(|E|)$

populate()

#3. (CONT'D) Sorting and filling the col_idx_weight array

PARALLEL VERSION

```
initLocks();
#pragma omp parallel {
    std::tuple<uint64_t, uint64_t, double> curr_edge;
    uint64_t curr_vertex, new_pos;

    #pragma omp for
    for (uint64_t i = 0; i < num_edges; i++) {
        curr_edge = e_list[i]; curr_vertex = std::get<0>(curr_edge);
        // locking on curr_vertex
        omp_set_lock(&count_locks[curr_vertex]);
        new_pos = count[curr_vertex] - 1; count[curr_vertex]--;
        // unlocking curr_vertex
        omp_unset_lock(&count_locks[curr_vertex]);

        col_idx_weight[new_pos] = std::make_pair(std::get<1>(curr_edge),
        std::get<2>(curr_edge));
    }

    // destroying locks and sorting neighbors
    #pragma omp parallel for
    for (uint64_t i = 0; i < num_vertices; i++){
        std::sort(col_idx_weight + row_ptr[i], col_idx_weight +
        row_ptr[i+1]);
        omp_destroy_lock(&count_locks[i]);
    }
}
```

Problem: in the parallel version of the counting sort the new_pos of a neighbor of curr_vertex depends on the order in which the threads access the edge list. Thus, BFS and DFS sums are different for each iteration.

Solution: sort the neighbors of every vertex in a predefined order (e.g. ascending order)

populate()

#3. (CONT'D) Sorting and filling the col_idx_weight array

PARALLEL VERSION

```
initLocks();
#pragma omp parallel {
    std::tuple<uint64_t, uint64_t, double> curr_edge;
    uint64_t curr_vertex, new_pos;

    #pragma omp for
    for (uint64_t i = 0; i < num_edges; i++) {
        curr_edge = e_list[i]; curr_vertex = std::get<0>(curr_edge);
        // locking on curr_vertex
        omp_set_lock(&count_locks[curr_vertex]);
        new_pos = count[curr_vertex] - 1; count[curr_vertex]--;
        // unlocking curr_vertex
        omp_unset_lock(&count_locks[curr_vertex]);

        col_idx_weight[new_pos] = std::make_pair(std::get<1>(curr_edge),
        std::get<2>(curr_edge));
    }

    // destroying locks and sorting neighbors
    #pragma omp parallel for
    for (uint64_t i = 0; i < num_vertices; i++){
        std::sort(col_idx_weight + row_ptr[i], col_idx_weight +
        row_ptr[i+1]);
        omp_destroy_lock(&count_locks[i]);
    }
}
```

Time Complexity:

$$O\left(\frac{|E|}{P} + \frac{|V|}{P} |V| \log(|V|)\right)$$


But the sort is done on an array that contains the neighbors of a single vertex. In a **sparse graph** we have that the number of neighbors of a vertex is $\ll |V|$. Therefore,

$$O\left(\frac{|E|}{P} + \frac{|V|}{P}\right)$$

where P is the number of partitions (i.e. threads)

populate()

Total time complexity

PARALLEL VERSION

$$\begin{aligned} & \mathcal{O}\left(\frac{|E|}{P} + P\right) + \mathcal{O}(|V|) + \mathcal{O}\left(\frac{|E|}{P} + \frac{|V|}{P}\right) \\ &= \mathcal{O}\left(\frac{|E|}{P} + P + |V| + \frac{|V|}{P}\right) \end{aligned}$$

With $|E| \gg |V|$: $\mathcal{O}\left(\frac{|E|}{P} + P\right)$

SERIAL VERSION

$$\begin{aligned} & \mathcal{O}(|E|) + \mathcal{O}(|V|) + \mathcal{O}(|E|) \\ &= \mathcal{O}(|V| + |E|) \end{aligned}$$

With $|E| \gg |V|$: $\mathcal{O}(|E|)$

get_neighbors()

```
CSRIter get_neighbors(uint64_t vertex_idx){
    return CSRIter(col_idx_weight + row_ptr[vertex_idx], col_idx_weight + row_ptr[vertex_idx + 1]);
}

class CSRIter {
    class iterator {
    public:
        iterator(std::pair<uint64_t, double> *ptr) : ptr(ptr) {}
        iterator operator++(){ ++ptr; return *this; }
        bool operator!=(const iterator &other) { return ptr != other.ptr; }
        const std::pair<uint64_t, double> &operator*(){return *ptr;};

    private:
        std::pair<uint64_t, double> *ptr;
    };

    private:
        std::pair<uint64_t, double> *begin_ptr;
        std::pair<uint64_t, double> *end_ptr;

    public:
        CSRIter(std::pair<uint64_t, double> *begin_ptr, std::pair<uint64_t, double> *end_ptr) :
        begin_ptr(begin_ptr), end_ptr(end_ptr) {}

        iterator begin() const { return iterator(begin_ptr); }

        iterator end() const { return iterator(end_ptr); }
    };
```

Spatial locality: neighbors of closer vertices are stored continuously

Comparison between serial CSR and the Adjacency List implementations

Populate CSR

x6.5 faster *

Populate AL

BFS - DFS CSR

x5.8 faster *

BFS - DFS AL

Graph: Wiki_talks (~ 5M edges, ~ 2.5M nodes)

*runned on system with: Intel Core i7-8550U with 1.8GHz clock speed, 4 cores, 8 Threads. The machine has 8GB of RAM, 256KB of L1 cache, 1M of L2 cache, and 8MB of L3 cache.

Comparison between Parallel and Serial Versions

Populate CSR - Parallel

x2.5 faster *




Populate CSR - Serial



Graph: Wiki_talks (~ 5M edges, ~ 2.5M nodes)

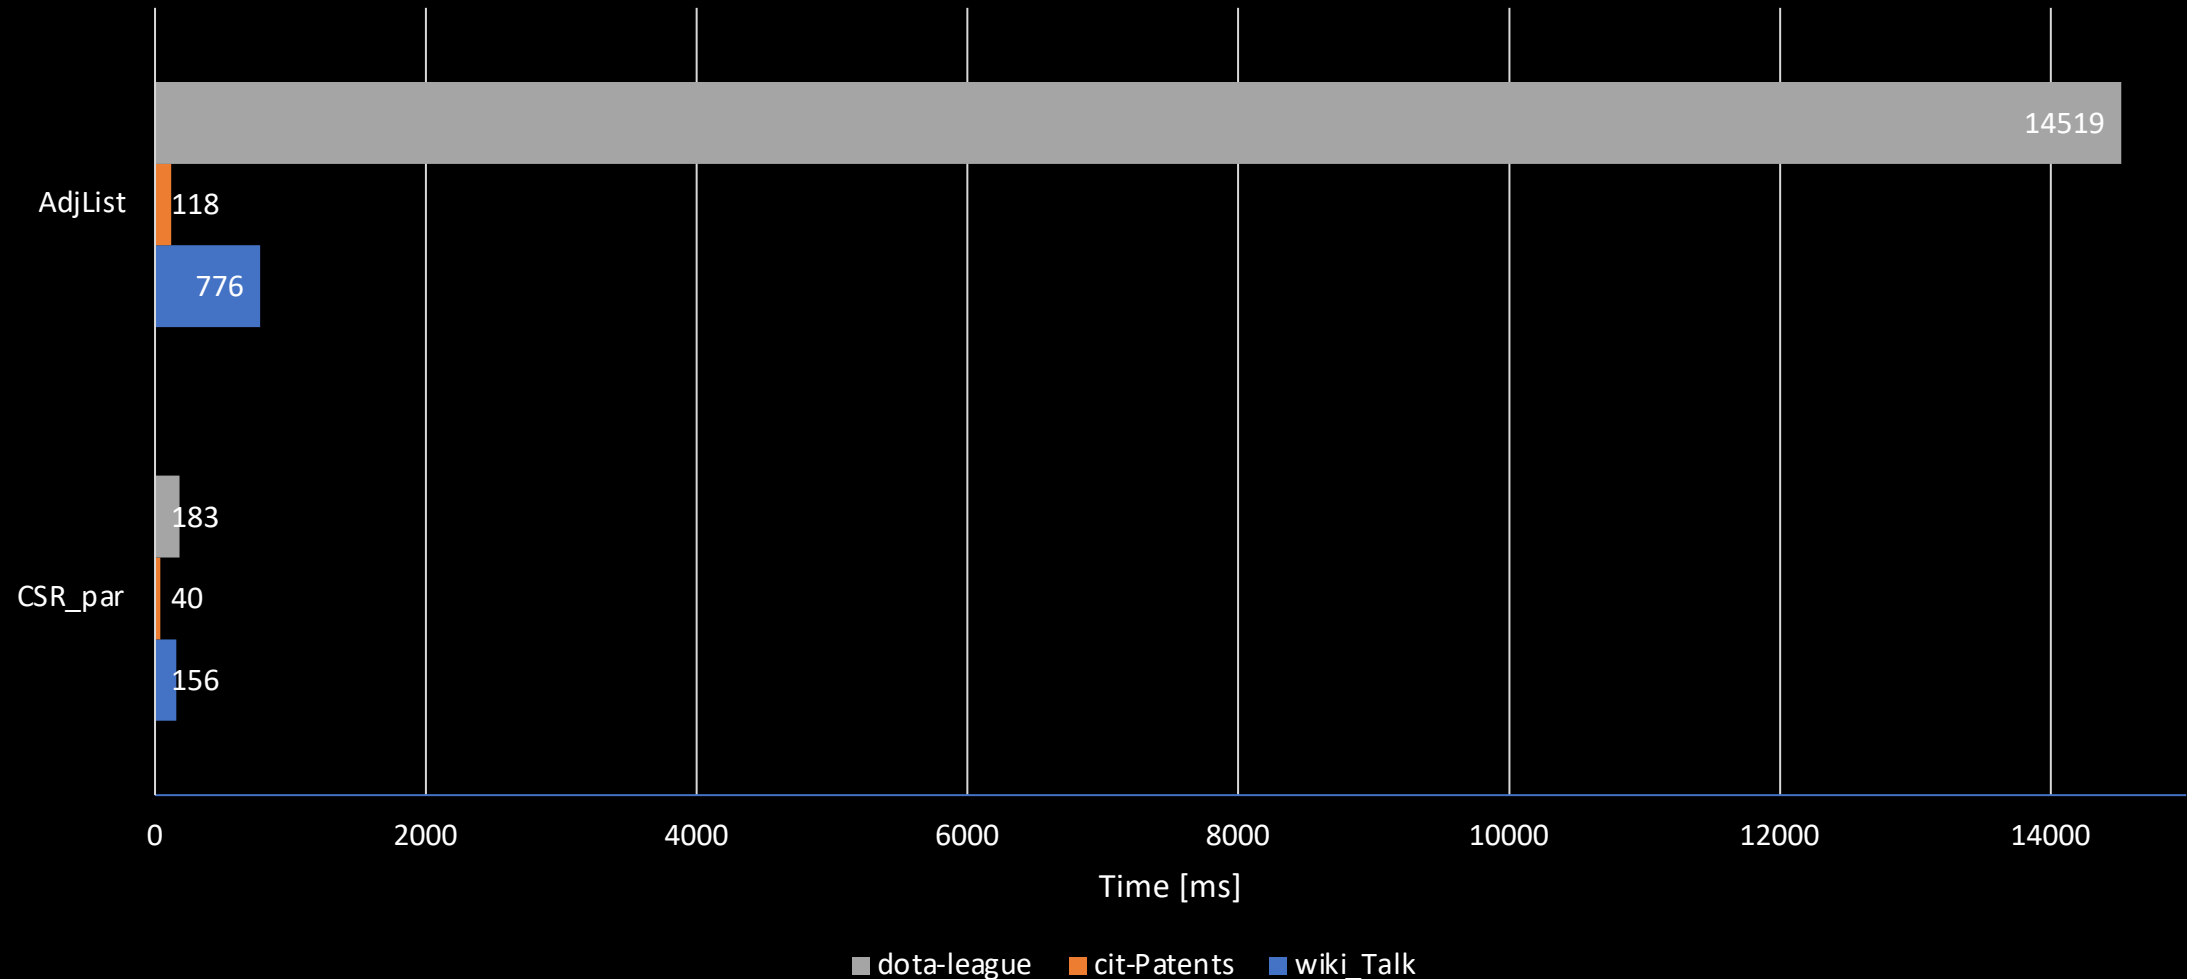
*runned on system with: Intel Core i7-8550U with 1.8GHz clock speed, 4 cores, 8 Threads. The machine has 8GB of RAM, 256KB of L1 cache, 1M of L2 cache, and 8MB of L3 cache.

Populating Time Results




■ dota-league ■ cit-Patents ■ wiki-Talk

BFS Time Results



System: nags33 @NECTSLab, 40 Threads

DFS Time Results



System: nags33 @NECTSLab, 40 Threads

Memory results


dota-league - AL (3640MB)



dota-league - CSR_par (1555MB)



cit-Patents - AL (1182MB)



cit-Patents - CSR_par (72MB)



wiki_Talk - AL (455MB)



wiki_Talk - CSR_par (9,5MB)



THANKS